

Module 2 Python Exercises

July 12, 2021

1 Module 2 Python Exercises: KNN and Perceptron

1.1 Filipp Krasovsky, 7-12-2021

```
[30]: import seaborn as sns
import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import matplotlib.pylab as pylab
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.preprocessing import OrdinalEncoder
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import OneHotEncoder
```

1.2 Drum Sounds Data

```
[91]: audio_data = pd.read_csv('audio_data.csv', index_col=0).
      ↳ drop(['label', 'filename'], axis=1)

X = audio_data.drop('label_text', axis=1)

le = preprocessing.LabelEncoder()
labels = audio_data['label_text']
le.fit(labels)
y = le.transform(labels)

u_labels = le.classes_

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
      ↳ random_state=42)
```

```

[92]: #train a knn model across k = 1 to 9 and get the manhattan and euclidean
      ↪performance(s).
k_val = []
euclidean_test = []
euclidean_train= []
manhattan_test = []
manhattan_train= []

for k in range(1,10):
    if k%2==0:
        continue
    k_val.append(k)
    #train using euclidean distance
    this_model = KNeighborsClassifier(n_neighbors=k,metric="euclidean")
    this_model = this_model.fit(X_train,y_train)
    #get training accuracy
    this_train = this_model.score(X_train,y_train)
    #get testing accuracy
    this_pred = this_model.predict(X_test)
    test_pred = accuracy_score(this_pred,y_test)
    #append
    euclidean_test.append(test_pred)
    euclidean_train.append(this_train)

    #train using manhattan distance
    this_model_manhattan =
    ↪KNeighborsClassifier(n_neighbors=k,metric="manhattan")
    this_model_manhattan = this_model_manhattan.fit(X_train,y_train)

    #get training accuracy
    manhattan_train_accuracy = this_model_manhattan.score(X_train,y_train)

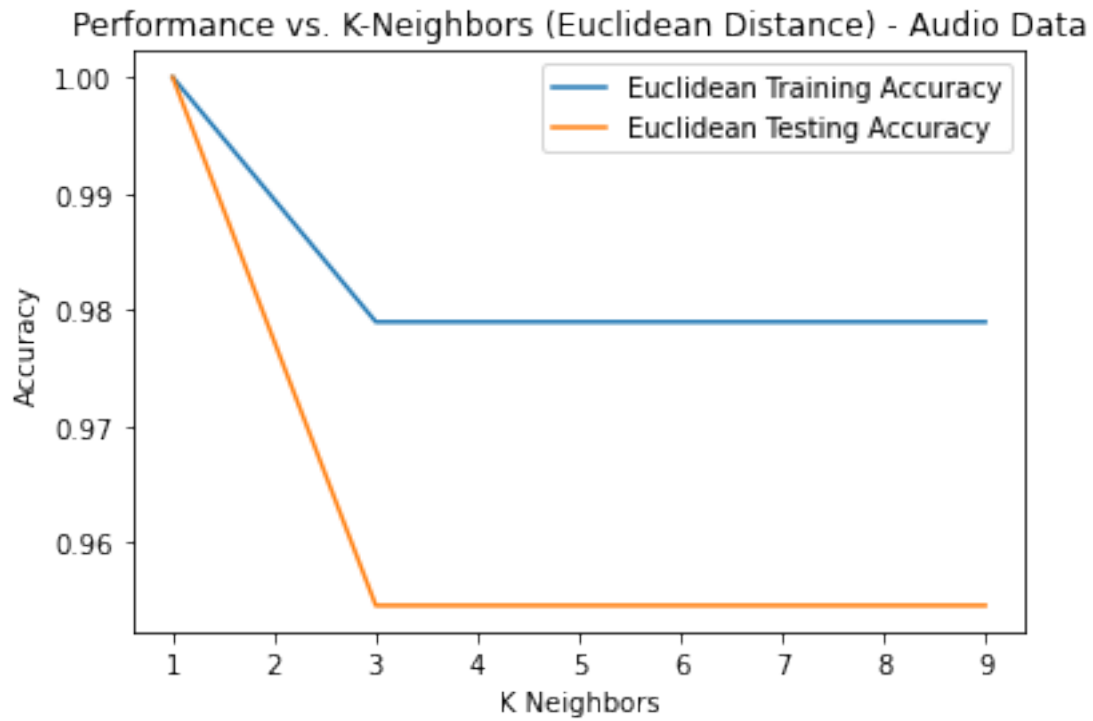
    #get testing accuracy
    this_pred_manhattan = this_model_manhattan.predict(X_test)
    this_pred_manhattan = accuracy_score(this_pred_manhattan,y_test)
    #append
    manhattan_test.append(this_pred_manhattan)
    manhattan_train.append(manhattan_train_accuracy)

```

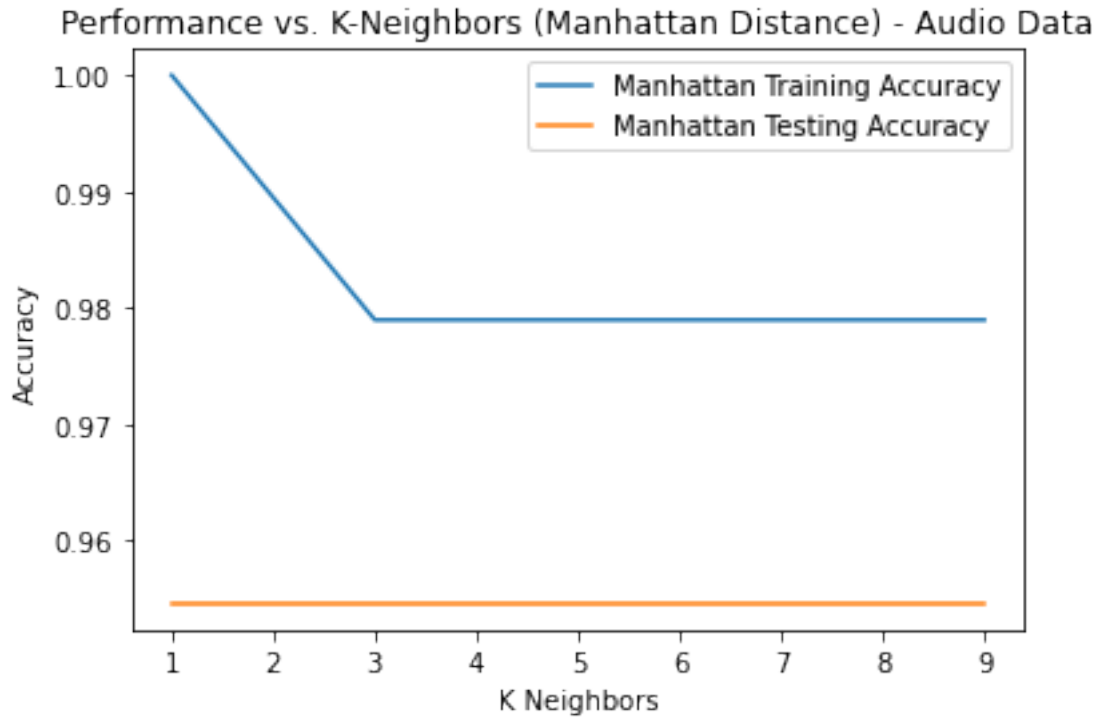
```

[93]: #plot euclidean test/train error
plt.plot(k_val, euclidean_train, label = "Euclidean Training Accuracy")
plt.plot(k_val, euclidean_test, label = "Euclidean Testing Accuracy")
plt.legend()
plt.xlabel("K Neighbors")
plt.ylabel("Accuracy")
plt.title("Performance vs. K-Neighbors (Euclidean Distance) - Audio Data")
plt.show()

```



```
[94]: #plot manhattan test/train error
plt.plot(k_val, manhattan_train, label = "Manhattan Training Accuracy")
plt.plot(k_val, manhattan_test, label = "Manhattan Testing Accuracy")
plt.legend()
plt.xlabel("K Neighbors")
plt.ylabel("Accuracy")
plt.title("Performance vs. K-Neighbors (Manhattan Distance) - Audio Data")
plt.show()
```



1.3 Animal Shelter Data

```
[27]: shelter_data = pd.read_csv('shelter_data.csv')
shelter_data.head()
```

```
[27]: AnimalID      Name      DateTime      OutcomeType OutcomeSubtype \
0  A671945  Hambone  2014-02-12 18:22:00  Return_to_owner      NaN
1  A656520   Emily  2013-10-13 12:44:00    Euthanasia    Suffering
2  A686464  Pearce  2015-01-31 12:28:00    Adoption    Foster
3  A683430    NaN  2014-07-11 19:09:00    Transfer    Partner
4  A667013    NaN  2013-11-15 12:52:00    Transfer    Partner
```

```
AnimalType SexuponOutcome AgeuponOutcome      Breed \
0      Dog  Neutered Male      1 year  Shetland Sheepdog Mix
1      Cat  Spayed Female      1 year  Domestic Shorthair Mix
2      Dog  Neutered Male      2 years      Pit Bull Mix
3      Cat  Intact Male      3 weeks  Domestic Shorthair Mix
4      Dog  Neutered Male      2 years  Lhasa Apso/Miniature Poodle
```

```
Color
0  Brown/White
1  Cream Tabby
2  Blue/White
```

3 Blue Cream
4 Tan

```
[110]: # this line drops any rows with missing data
cleaned_shelter_data = shelter_data.dropna()

# here we grab the data we want from pandas
X_data = cleaned_shelter_data[['AnimalType', 'SexuponOutcome', 'AgeuponOutcome']]
y_data = cleaned_shelter_data[['OutcomeType']]

#use the oneHotEncoder to transform our variables and get label names
ohe = OneHotEncoder(sparse=False)
feature_array = ohe.fit_transform(X_data).toarray()
feature_labels= ohe.categories_
#convert labels into one array
feature_labels = np.concatenate(feature_labels)
#combine labels and data
features = pd.DataFrame(feature_array, columns = feature_labels)

le = preprocessing.LabelEncoder()
le.fit(y_data)
y = le.transform(y_data)

X_train, X_test, y_train, y_test = train_test_split(features, y, test_size=0.
↪10, random_state=42)
```

C:\Users\13234\Miniconda3\lib\site-packages\sklearn\utils\validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
return f(*args, **kwargs)

```
[111]: #train a knn model across k = 1 to 9 and get the cosine performance for
↪training and testing.
k_val = []
cosine_train = []
cosine_test = []
for k in range(1,10):
    if k%2==0:
        continue
    k_val.append(k)
    #train using cosine distance
    this_model = KNeighborsClassifier(n_neighbors=k,metric="cosine")
    this_model = this_model.fit(X_train,y_train)
    #get training accuracy
    this_train = this_model.score(X_train,y_train)
    #get testing accuracy
```

```

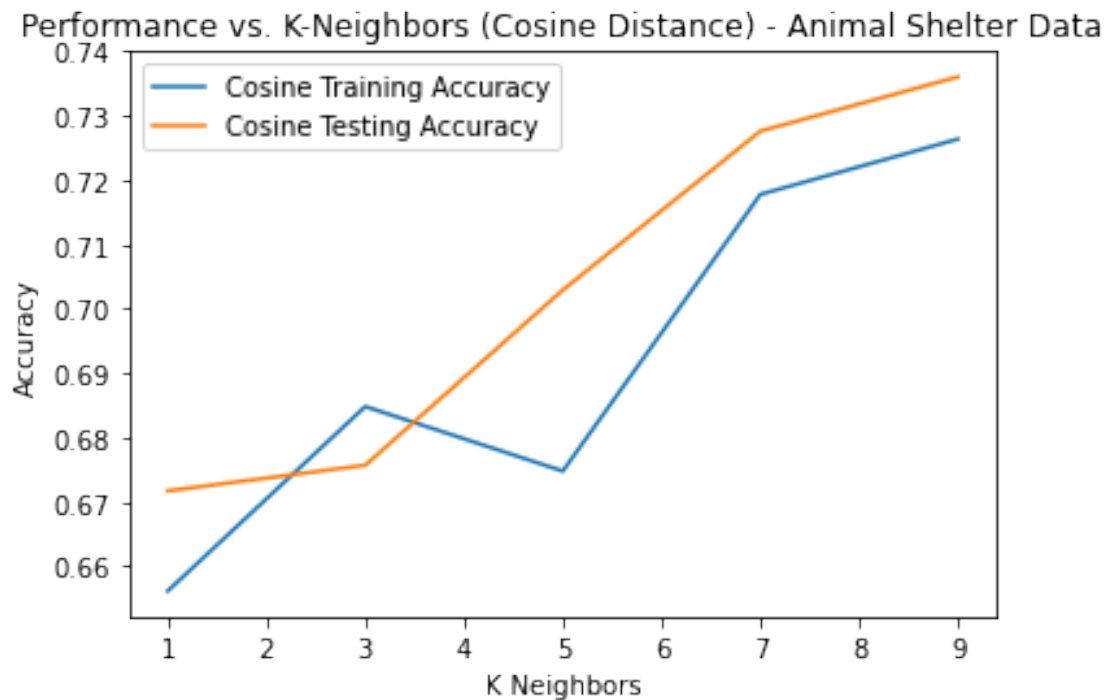
this_pred = this_model.predict(X_test)
test_pred = accuracy_score(this_pred,y_test)
#append
cosine_train.append(test_pred)
cosine_test.append(this_train)

```

```

[112]: #plot cosine test/train error
plt.plot(k_val, cosine_train, label = "Cosine Training Accuracy")
plt.plot(k_val, cosine_test, label = "Cosine Testing Accuracy")
plt.legend()
plt.xlabel("K Neighbors")
plt.ylabel("Accuracy")
plt.title("Performance vs. K-Neighbors (Cosine Distance) - Animal Shelter Data")
plt.show()

```



2 Text Data

```

[113]: text_data = pd.read_csv('text_data.csv',index_col=0).drop('meta_title',axis=1)
display(text_data)

```

	meta_author	000	10	11	13	136	13th	1648	1683	1685	...	yielding	\
0	hamilton	0	0	0	0	0	0	0	0	0	...	0	
1	jay	0	0	0	0	0	0	0	0	0	...	0	
2	jay	0	0	0	0	0	0	0	0	1	...	0	

3	jay	0	0	0	0	0	0	0	0	0	...	0
4	jay	0	0	0	0	0	0	0	0	0	...	0
..
80	hamilton	0	0	0	0	0	0	0	0	0	...	0
81	hamilton	0	0	0	0	0	0	0	0	0	...	0
82	hamilton	0	0	0	0	0	0	0	0	0	...	0
83	hamilton	0	0	0	0	1	0	0	0	0	...	0
84	hamilton	0	0	0	0	0	0	0	0	0	...	0

	yoke	yokes	york	young	yourselves	zaleucus	zeal	zealand	zealous
0	0	0	1	0	0	0	3	0	0
1	0	0	1	0	0	0	0	0	0
2	0	0	1	0	0	0	0	0	0
3	0	0	1	0	0	0	0	0	0
4	0	0	1	1	1	0	0	0	0
..
80	0	0	2	0	0	0	0	0	0
81	0	0	2	0	0	0	0	0	0
82	0	0	5	0	0	0	0	0	0
83	0	0	3	0	0	0	2	0	0
84	0	0	1	0	0	0	1	0	2

[85 rows x 8561 columns]

```
[115]: #transform the output label
X = text_data.drop('meta_author',axis=1)
le = preprocessing.LabelEncoder()
labels = text_data['meta_author']
le.fit(labels)
y=le.transform(labels)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
↪random_state=42)

[116]: #from k = 1 to 9 (odd) train a euclidean and cosine distance for performance.
euclidean_train = []
euclidean_test  = []
cosine_train    = []
cosine_test     = []
k_val = []

for k in range(1,10):
    if k%2==0:
        continue
    k_val.append(k)
    #train using cosine distance
    this_model = KNeighborsClassifier(n_neighbors=k,metric="cosine")
```

```

this_model = this_model.fit(X_train,y_train)
#get training accuracy
this_train = this_model.score(X_train,y_train)
#get testing accuracy
this_pred = this_model.predict(X_test)
test_pred = accuracy_score(this_pred,y_test)
#append
cosine_train.append(test_pred)
cosine_test.append(this_train)


#train using euclidean distance
this_model = KNeighborsClassifier(n_neighbors=k,metric="euclidean")
this_model = this_model.fit(X_train,y_train)
#get training accuracy
this_train = this_model.score(X_train,y_train)
#get testing accuracy
this_pred = this_model.predict(X_test)
test_pred = accuracy_score(this_pred,y_test)
#append
euclidean_train.append(test_pred)
euclidean_test.append(this_train)

```

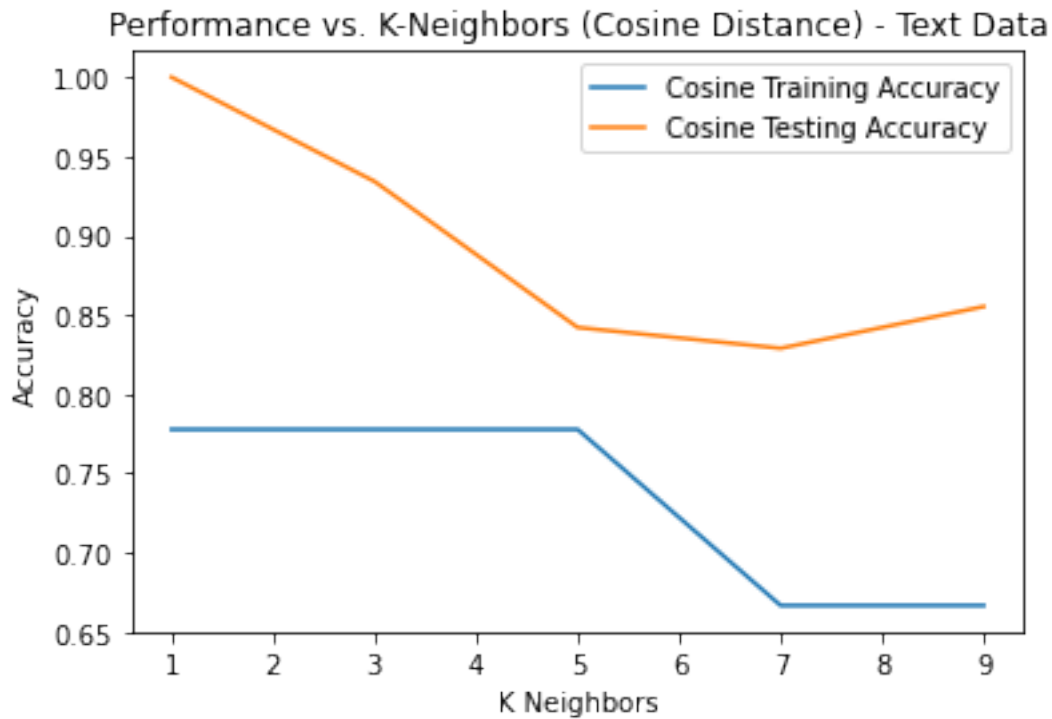
```

[117]: #plot euclidean performance for text data
plt.plot(k_val, euclidean_train, label = "Euclidean Training Accuracy")
plt.plot(k_val, euclidean_test, label = "Euclidean Testing Accuracy")
plt.legend()
plt.xlabel("K Neighbors")
plt.ylabel("Accuracy")
plt.title("Performance vs. K-Neighbors (Euclidean Distance) - Text Data")
plt.show()

```

```
[118]: #plot cosine performance for text data
plt.plot(k_val, cosine_train, label = "Cosine Training Accuracy")
plt.plot(k_val, cosine_test, label = "Cosine Testing Accuracy")
plt.legend()
plt.xlabel("K Neighbors")
plt.ylabel("Accuracy")
plt.title("Performance vs. K-Neighbors (Cosine Distance) - Text Data")
plt.show()
```



3 Perceptron Implementation

begin by reloading our audio data.

```
[174]: audio_data = pd.read_csv('audio_data.csv', index_col=0).
        ↳ drop(['label', 'filename'], axis=1)
X = audio_data.drop('label_text', axis=1)
le = preprocessing.LabelEncoder()
labels = audio_data['label_text']
le.fit(labels)

#for the perceptron, we transform our 0 and 1 classes into -1 and 1.
y=le.transform(labels)
y[y==0] = -1

u_labels = le.classes_
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
        ↳ random_state=42)
```

```
[183]: #create a sign function for the training and testing stages.
```

```
def sign(dp):
    if (dp!=0):
        return(dp/abs(dp))
    else:
```

```
return(0)
```

```
[368]: #take in a set of predictors and their weight vectors and return the sign of
        ↳ their dot product.
        #iterate over each object in the test set, multiply it by the supplied weight
        ↳ vector, and append it to a prediction vector
        #after operating over it with the sign function.
```

```
def predict_perceptron(X,w):
    predictions=[]
    for i in range(0,X.shape[0]):
        prediction = sign(np.dot(X.iloc[i].ravel(),w))
        predictions.append(prediction)
    return(predictions)
```

```
[351]: #assumption - we require the response variable to be stored in a separate
        ↳ vector y
```

```
def train_perceptron(X,y):
    #initialize a weight vector of size D (number of dimensions in X)
    D = X.shape[1]
    R = X.shape[0]
    w = np.array(np.zeros(D))
    x = np.array(np.zeros(D))

    #arbitrary max iteration
    maxIt = 100

    #iterate over each training sample. if the dot product sign != output sign,
    ↳ update our weight vector.
    for i in range(1,maxIt):
        for j in range(0,R):
            this_row = np.array(X.iloc[[j]])
            this_y = y[j]
            dp = np.dot(this_row.ravel(),w)
            a = dp

            if(a*this_y<=0):
                w = w + (this_row.ravel() * this_y)

    return(w)
```

```
[372]: #train our perceptron model and then predict for our test set.
        perceptron = train_perceptron(X_train,y_train)
        p_predict = predict_perceptron(X_test,perceptron)
        p_accuracy = accuracy_score(p_predict,y_test)
        print("Perceptron Accuracy: %",p_accuracy)
```

Perceptron Accuracy: % 0.9545454545454546

```
[374]: #verify accuracy by looking at the scikit perceptron  
from sklearn.linear_model import Perceptron  
sk_perceptron = Perceptron()  
sk_perceptron.fit(X_train,y_train)  
sk_predict = sk_perceptron.predict(X_test)  
print("SK Perceptron Accuracy: %",accuracy_score(sk_predict,y_test))
```

SK Perceptron Accuracy: % 1.0

Our Accuracies are within 4.5% of each other and both considerably high, so we can claim some level of success in recreating the perceptron by hand.