

# Untitled

August 3, 2021

## 1 Module 5 Homework Assignments

Neural Networks and Gradient Descent Filipp Krasovsky, August 1, 2021

### 1.1 Dataset: MNIST

The MNIST dataset is a set of images of handwritten digits from 0-9. They are a series of small images that have been aligned and scaled to be a similar size and orientation. I recommend loading the data using Keras if possible <https://keras.io/api/datasets/mnist/>.

```
[1]: from datetime import datetime
import tensorflow as tf
from tensorflow import keras
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import cross_val_score
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from sklearn.linear_model import Perceptron
from sklearn.metrics import plot_confusion_matrix, classification_report
from sklearn.neural_network import MLPClassifier
from mlxtend.plotting import plot_decision_regions
```


### 1.2 The XOR Problem

Using the provided sample code, create sample data for the XOR problem. Plot these data, color-coded by the labels (category). Use sklearn's perceptron function to train a classifier on these data. What error do you obtain? Plot the confusion matrix and use metrics.classification\_report to make a table of classification results.

```
[3]: #create sample XOR data.
rng = np.random.RandomState(0)
X = rng.randn(1000, 2)
y = np.logical_xor(X[:, 0] > 0, X[:, 1] > 0)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
↳random_state=42)
```

```
[4]: #plot data, color coded by label
cdict = {True: 'blue', False: 'orange'}


fig, ax = plt.subplots()
for g in np.unique(y_train):
    ix = np.where(y_train == g)
    ax.scatter(X_train[ix,0],X_train[ix,1],c = cdict[g], label = g, s = 5)
ax.legend()
plt.show()
```



```
[5]: #use the sklearn perceptron to train a classifier on this dataset and observe
      ↳ the error:
clf = Perceptron( random_state=42)
clf.fit(X_train,y_train)
clf.score(X_train,y_train)
clf_pred = clf.predict(X_test)
```

```
[6]: #plot confusion matrix and classification results
plot_confusion_matrix(clf,X_train,y_train)
```

```
[6]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2503f2be820>
```



```
[7]: print(classification_report(clf_pred,y_test))
```

	precision	recall	f1-score	support
False	0.00	0.00	0.00	0
True	1.00	0.32	0.48	100
accuracy			0.32	100
macro avg	0.50	0.16	0.24	100
weighted avg	1.00	0.32	0.48	100

```
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\metrics\_classification.py:1245: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\metrics\_classification.py:1245: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\metrics\_classification.py:1245: UndefinedMetricWarning: Recall
and F-score are ill-defined and being set to 0.0 in labels with no true samples.
Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

In conclusion, we get incredibly low accuracy and precision for the False class and high precision for the True class. Our perceptron classifier performs poorly.

```
[8]: #Next we create a neural net with backpropagation:
nnet = MLPClassifier(hidden_layer_sizes=(5,), max_iter=1000, alpha=1e-4,
solver='lbfgs', verbose=True, random_state=1,
learning_rate_init=.1)
```


```
[9]: #fit model
nnet.fit(X_train,y_train)
#predict
nnet_pred = nnet.predict(X_test)
```

```
[10]: print(classification_report(nnet_pred,y_test))
```

	precision	recall	f1-score	support
False	0.99	0.99	0.99	68
True	0.97	0.97	0.97	32
accuracy			0.98	100
macro avg	0.98	0.98	0.98	100
weighted avg	0.98	0.98	0.98	100

```
[11]: #plot the test set and our decision boundary.
plot_decision_regions(X_test,y_test.astype(np.int64),clf=nnet,legend=2)
```

```
[11]: <AxesSubplot:>
```



### 1.3 Moons Dataset


```
[12]: from sklearn.datasets import make_moons
X,y = make_moons(n_samples=100,noise=0.3, random_state=0)

[13]: #next, make a range of 1 to 20 hidden units and train the nnet with each value,
      ↪plotting as we go:
hidden_layers = range(1,21)

[14]: for hidden_layer in hidden_layers:
      #create our neural net
      nnet = MLPClassifier(hidden_layer_sizes=(hidden_layer,), max_iter=1000,
      ↪alpha=1e-4,
      solver='lbfgs', verbose=True, random_state=1,
      learning_rate_init=.1)
      #fit model
      nnet.fit(X,y)
      #plot the train set and our decision boundary.
      plot_decision_regions(X,y,clf=nnet,legend=2)
      plt.title('Neural Net on Moons Dataset, Hidden Layer Size: '+
      ↪str(hidden_layer))
      plt.show()
```


C:\Users\13234\Miniconda3\lib\site-packages\mlxtend\plotting\decision\_regions.py:244: UserWarning: No contour levels were found within the data range.

```
ax.contour(xx, yy, Z, cset.levels,
```






```
C:\Users\13234\Miniconda3\lib\site-  
packages\mlxtend\plotting\decision_regions.py:244: UserWarning: No contour  
levels were found within the data range.  
    ax.contour(xx, yy, Z, cset.levels,
```

Neural Net on Moons Dataset, Hidden Layer Size: 2







Neural Net on Moons Dataset, Hidden Layer Size: 3














Neural Net on Moons Dataset, Hidden Layer Size: 10




Neural Net on Moons Dataset, Hidden Layer Size: 11




Neural Net on Moons Dataset, Hidden Layer Size: 12




Neural Net on Moons Dataset, Hidden Layer Size: 13




Neural Net on Moons Dataset, Hidden Layer Size: 14




Neural Net on Moons Dataset, Hidden Layer Size: 15




Neural Net on Moons Dataset, Hidden Layer Size: 16




Neural Net on Moons Dataset, Hidden Layer Size: 17




Neural Net on Moons Dataset, Hidden Layer Size: 18



Neural Net on Moons Dataset, Hidden Layer Size: 19





#### 1.4 MNIST Dataset

```
[25]: #load in our image dataset from keras.
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
#sanity check
assert x_train.shape == (60000, 28, 28)
assert x_test.shape == (10000, 28, 28)
assert y_train.shape == (60000,)
assert y_test.shape == (10000,)

#reshape our data to be compatible with sklearn.
nsamples, nx, ny = x_train.shape
x_train_d2 = x_train.reshape((nsamples,nx*ny))
nsamples, nx, ny = x_test.shape
x_test_d2 = x_test.reshape((nsamples,nx*ny))

[26]: #use the sklearn perceptron to train a classifier on this dataset and observe
      ↳ the error:
clf = Perceptron( random_state=42)
clf.fit(x_train_d2,y_train)
clf_pred = clf.predict(x_test_d2)
accuracy_score(clf_pred,y_test)
```

[26]: 0.8892



```
[36]: #optimize across three different metrics:
hidden_units = [1,5,10,20]
solver = ['adam','sgd','lbfgs']
activation = ['relu','logistic']
scores = {}

for a in activation:
    scores[a] = {}
    for s in solver:
        scores[a][s] = []
        for h in hidden_units:
            #train our classifier
            clf = MLPClassifier(
                hidden_layer_sizes=(h,),
                max_iter=1000,
                alpha=1e-4,
                solver=s,
                activation=a,
                random_state=1,
                learning_rate_init=.1)
            #fit, record accuracy
            clf.fit(x_train_d2,y_train)
            clf_pred = clf.predict(x_test_d2)
            scores[a][s].append(accuracy_score(clf_pred,y_test))
        plt.plot(hidden_units,scores[a][s],label=a+"/"+s)
plt.legend()
plt.xlabel("Hidden Units")
plt.ylabel("Accuracy")
plt.title("MNIST Data Classification Performance")
plt.show()
```

```
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:614:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000) reached and
the optimization hasn't converged yet.
    warnings.warn(
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:500:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

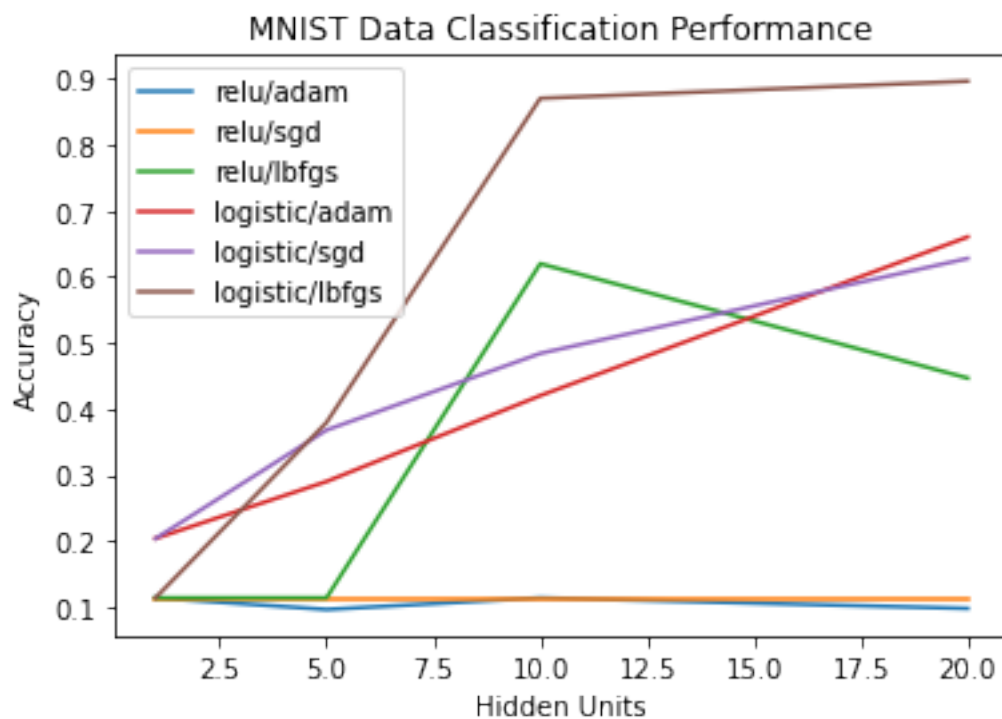
```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
C:\Users\13234\Miniconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:500:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
    <https://scikit-learn.org/stable/modules/preprocessing.html>  
    self.n\_iter\_ = \_check\_optimize\_result("lbfgs", opt\_res, self.max\_iter)  
C:\Users\13234\Miniconda3\lib\site-  
packages\sklearn\normal\_distribution\multilayer\_perceptron.py:500:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
    <https://scikit-learn.org/stable/modules/preprocessing.html>  
    self.n\_iter\_ = \_check\_optimize\_result("lbfgs", opt\_res, self.max\_iter)  
C:\Users\13234\Miniconda3\lib\site-  
packages\sklearn\normal\_distribution\multilayer\_perceptron.py:500:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
    <https://scikit-learn.org/stable/modules/preprocessing.html>  
    self.n\_iter\_ = \_check\_optimize\_result("lbfgs", opt\_res, self.max\_iter)  
C:\Users\13234\Miniconda3\lib\site-  
packages\sklearn\normal\_distribution\multilayer\_perceptron.py:500:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
    <https://scikit-learn.org/stable/modules/preprocessing.html>  
    self.n\_iter\_ = \_check\_optimize\_result("lbfgs", opt\_res, self.max\_iter)



[37]: scores

```
[37]: {'relu': {'adam': [0.1135, 0.0958, 0.1135, 0.0974],
  'sgd': [0.1135, 0.1135, 0.1135, 0.1135],
  'lbfgs': [0.1135, 0.1135, 0.62, 0.4466]},
  'logistic': {'adam': [0.2035, 0.2898, 0.42, 0.6606],
  'sgd': [0.2033, 0.3678, 0.484, 0.628],
  'lbfgs': [0.1135, 0.3793, 0.8707, 0.8967]}}
```

[ ]: